

Population Growth Analysis and Forecasting Using Data Analytics

by

Nischitha Purushotham

University of Europe for Applied Sciences

Contents

1. Introduction
2. Objectives
3. Dataset Description
4. Data Source and References
5. Data Cleaning Process
6. Methodology
7. Analysis and Findings
8. Charts and Visualisations
9. Forecasting Results
10. Interpretation of Key Insights
11. Limitations of the Analysis
12. Conclusion
13. Appendices
14. References

Abstract

This report presents a population analytics study built from the World Bank population dataset and processed in a Jupyter Notebook using Python, pandas, matplotlib, seaborn, and scikit-learn. The project transforms a wide annual population file into a country-year panel, cleans the data, creates analytical variables such as annual change, growth rate, decade grouping, and continent mapping, and then applies exploratory analysis, visualisation, and forecasting. The cleaned dataset covers a long historical window from 1960 to 2024 and includes 265 countries and regions. The analysis shows a strong concentration of population in a small number of countries, especially India, China, the United States, Indonesia, Pakistan, Nigeria, and Brazil. A comparison of growth patterns reveals clear differences across countries, continents, and decades, while the linear regression forecast indicates that India is likely to continue growing through 2030. The report also

discusses the limitations of using historical population series for prediction, especially when broader demographic, policy, and migration forces are not included.

1. Introduction

Population growth is one of the most important indicators used in demographic, economic, and social analysis. It affects healthcare systems, education planning, housing demand, transportation networks, labour markets, food supply, and environmental sustainability. When population changes are tracked over time, policy makers and researchers can better understand how societies develop, where pressure on infrastructure is likely to increase, and which regions may need targeted support.

This project applies data analytics methods to historical population data to explore population change across countries and regions. The analysis is carried out in Python using a structured workflow that starts with raw data inspection, continues with cleaning and feature creation, and then moves into exploratory analysis, visualisation, and forecasting. The aim is not only to describe what happened historically, but also to use simple predictive modelling to estimate how the population of selected countries may evolve in the near future.

1.1 Project Context

The notebook begins with a World Bank population file that is originally arranged in a wide format, with years stored as separate columns. The file also contains a metadata companion dataset that is used to associate countries with broader geographic regions. After reshaping and cleaning, the dataset becomes suitable for year-by-year population trend analysis, continent-level comparison, and regression-based forecasting.

2. Objectives

- Collect a long-run population dataset and prepare it for analysis.
- Clean the raw data by handling formatting issues, missing values, and duplicate or non-country entries.
- Create useful analytical fields such as annual change, percentage growth rate, decade, growth category, and continent.
- Identify countries with the highest and lowest population levels in the latest year.
- Compare population growth patterns across selected countries and continents.
- Visualise the results using line charts, bar charts, scatter plots, pie charts, and heatmaps.

- Use linear regression to forecast population values for India and for the overall dataset trend.

3. Dataset Description

The analysis is based on the World Bank indicator file for Population, total (SP.POP.TOTL). In the raw form used in the notebook, the dataset contains 266 rows and 71 columns, with each row representing a country, region, or aggregate entity, and with yearly population values stored across columns from 1960 to 2024. The notebook then reshapes the data into a long format so that each row represents a country-year observation. This structure is more useful for time-series analysis, grouping, and forecasting.

The notebook output shows 265 unique entities after grouping by country name, which indicates that the source includes both individual countries and some broader entities. The study period covers 65 years, which is far more than the minimum required 30 years in the project brief. This makes the dataset suitable for analysing long-run demographic change, decade-level patterns, and recent short-term forecasts.

World Bank: <https://data.worldbank.org/indicator/SP.POP.TOTL>

3.1 Dataset Summary

Metric	Value
Source	World Bank Population, total dataset
Raw structure	Wide format, yearly columns
Raw dimensions	266 rows x 71 columns
Unique entities	265 countries/regions
Time span	1960 to 2024
Analytical format	Long country-year panel

4. Data Source and References

The primary source used in the notebook is the World Bank population database. The raw CSV file was imported directly into pandas, and the metadata file was used to map country codes to regions. This combination is especially useful because it allows the analysis to move beyond simple country-level totals and also compare broad geographic patterns. The notebook filters the dataset to keep the Population, total indicator and then removes aggregate or non-country entries where required for cleaner visualisation and forecasting.

Using the World Bank as the core source is appropriate for an academic project because the data is consistent, widely recognised, and regularly maintained. It also enables the use of a single analytical framework across many countries and over many decades. The metadata file adds a second layer of value by allowing each country to be associated with a continent or region, which is essential for grouped analysis and heatmap creation.

5. Data Cleaning Process

5.1 Dataset Inspection

The first step was to inspect the raw file structure using `df.info()` and `df.columns`. This confirmed that the source was originally in a wide layout with one row per entity and separate columns for each year. The notebook also revealed an extra unnamed column at the end, which is a common issue in downloaded CSV files and requires careful handling.

5.2 Reshaping the Data

The year columns were transformed into two variables, Year and Population, using a melt operation. This converted the dataset into a long format where each row contains a country, a year, and the corresponding population value. The reshaped format is much easier to use for trend analysis, grouping, and forecasting.

5.3 Handling Missing Values

Missing values were checked and removed where they interfered with the analysis. The notebook used `dropna()` in several places to clean incomplete records. Because some missing values were tied to aggregate or special entries rather than genuine country records, the filtering stage later removed many of them more effectively than simple deletion alone.

5.4 Duplicate Check

Duplicate records were checked and removed with `drop_duplicates()` when necessary. This is an important quality control step because repeated rows can distort summary statistics, heatmaps, and forecast models. The cleaned working file was saved as `cleaned_population.csv`.

5.5 Data Type Conversion

The year was converted to an integer format and the population to a numeric format. This ensures that time-based grouping, arithmetic, and regression modelling work correctly. Without this conversion, the notebook would not be able to calculate annual change or fit a linear model.

5.6 Column Renaming

The notebook renamed Country Name to Country so that the column names were shorter, clearer, and easier to use in later analysis. Keeping readable names also improves the final report and reduces the chance of coding errors.

5.7 Feature Engineering

Several useful analytical columns were created. Annual Change was calculated with the difference between consecutive years for each country. Growth Rate was calculated with percentage change. Decade was derived from the year value. A continent field was created by merging the metadata region mapping.

5.8 Filtering Non-Country Entities

The notebook removed rows containing World, income groups, IBRD, IDA, and other aggregate labels when those values were not relevant to the country-level analysis. It also checked that country codes were three-letter alphabetic codes and retained only the SP.POP.TOTL indicator. This step reduced noise and made the charts more interpretable.

6. Methodology

The project follows a standard data analytics workflow with five connected stages. Each stage builds on the previous one so that the final result is both technically consistent and easy to explain in a report.

- **Data Collection:** The notebook imports the World Bank CSV file and its metadata companion file. This provides the historical population values and the geographic context needed for broader comparisons.
- **Data Preparation:** The raw data is reshaped from wide to long form, types are corrected, missing and duplicate entries are handled, and several helpful features are added for later analysis.
- **Exploratory Data Analysis:** The cleaned data is summarised to identify top countries, basic trends, average values, and year-by-year changes. This stage answers the question of what the historical data is showing.
- **Advanced Analysis:** Heatmaps and comparison plots are used to study growth patterns across decades, regions, and country size. This stage helps reveal relationships that are not obvious in the basic tables.

- **Forecasting:** Linear regression is used to project future population values for India and the combined population trend for the dataset. The model is intentionally simple so that the results are easy to interpret in an academic setting.

7. Analysis and Findings

The cleaned dataset provides a broad view of population distribution across countries and regions. The most striking observation is that the population is concentrated heavily in a small number of countries. The notebook output for the latest year shows India, China, the United States, Indonesia, Pakistan, Nigeria, Brazil, Bangladesh, the Russian Federation, and Ethiopia as the top ten entities by population.

India and China sit far above the rest of the list, which is visible in both the table output and the bar chart. The gap between the first two countries and the remaining countries is substantial, showing that the world's largest population bases are highly concentrated in Asia. At the same time, the selected top countries include both very large developing economies and industrialised countries, which makes the comparison useful for understanding different demographic pathways.

The heatmap-based analysis shows that growth rates vary by continent and decade. Middle East and North Africa, South Asia, and Sub-Saharan Africa show stronger growth values in the earlier decades, while Europe and Central Asia remain comparatively lower across the same period. The decade pattern suggests that growth has generally moderated in several regions over time, even though population totals in many countries continue to rise.

The country-level growth heatmap makes it easy to spot countries with unusually high or negative growth in particular decades. This kind of chart is useful because it highlights both sustained trends and exceptional cases. Some countries show strong early acceleration, while others display flattening growth or a short-term decline in later years.

7.1 Comparative Reading of the Notebook Outputs

One useful way to read the notebook is to separate absolute population size from relative growth. Large countries such as India, China, and the United States naturally dominate absolute rankings because they have large base populations. However, the percentage growth charts show that smaller countries can still post meaningful changes over time. This is why the report uses both growth rate and total population measures instead of relying on only one view.

The comparison between annual change and growth rate is especially important. Annual change tells us how many people were added or lost in a given year, while growth rate tells us how fast that

change occurred relative to the size of the population. Together, these measures show whether a country is expanding because it is already large or because it is growing quickly from a smaller base.

The notebook also suggests that continental patterns are not static. Some regions exhibit slowing growth in later decades, which may reflect changes in fertility, urbanisation, education, and health conditions. Others remain relatively strong, indicating that demographic transition is progressing at different speeds in different parts of the world.

7.2 Top 10 Countries by Population in 2024

Rank	Country	Population
1	India	1.450936e+09
2	China	1.408975e+09
3	United States	3.401110e+08
4	Indonesia	2.834879e+08
5	Pakistan	2.512692e+08
6	Nigeria	2.326795e+08
7	Brazil	2.119986e+08
8	Bangladesh	1.735624e+08
9	Russian Federation	1.435339e+08
10	Ethiopia	1.320598e+08

8. Charts and Visualisations

The notebook includes a set of visualisations built in matplotlib and seaborn. Together, they show population concentration, growth structure, continent-level differences, and forecasting output. Each chart is accompanied by a brief interpretation below so that the report remains readable even without the notebook.




Fig. 1: Top 10 Countries by Population (2024)

The bar chart confirms that India and China are the clear population leaders in the latest year. The remaining countries still have very large populations, but the scale drops sharply after the first two entries. This indicates strong concentration in a small number of countries.

Top 5 Most Populated Countries




Fig. 2: Top 5 Most Populated Countries

The pie chart shows how much of the top-five group is made up by India and China. Their combined share dominates the chart, which highlights how concentrated the upper end of the population ranking is.

Bottom 10 Least Populated Countries




Fig.3: Bar Chart — Bottom 10 Least Populated Countries

Tuvalu is the least populated country with approximately 9,646 people. Several of these nations also show negative population growth due to emigration and low birth rates.

Correlation Heatmap




Fig. 5: Correlation Heatmap — Population, Growth Rate, Annual Change

Population size and Annual Change are strongly positively correlated — larger countries contribute more absolute population growth. Growth Rate shows a weaker correlation with population size, meaning smaller nations can still have high percentage growth rates.

Growth Rate Distribution — Histogram




Fig.6: Histogram — Distribution of Growth Rates Across All Countries

The majority of country-year observations fall in the 0–3% growth rate range. The distribution is right-skewed, with a long tail of extreme positive growth rates and a smaller cluster of negative growth rates.




Fig. 7: Growth Rate Heatmap

The country-level heatmap shows a wide spread of growth patterns across decades. Some countries have strong positive growth in earlier periods, while others flatten or turn negative in later years. This visual is useful for spotting long-run demographic change.




Fig. 8: Average Growth Rate by Continent and Decade

The continent-level heatmap shows that growth differs across both geography and time. The Middle East, North Africa and Sub-Saharan Africa generally display stronger growth than Europe and Central Asia. Several regions also show a clear slowing trend in the more recent decades.




Fig. 9: Population Forecast for India

The India forecast uses a linear regression trend line to extend the historical series into 2030. The predicted values continue the upward pattern visible in the actual data, suggesting that India will remain in a strong growth phase over the short term.




Fig. 10: Population Forecast for the Combined Dataset

The combined forecast indicates that the broader population trend in the dataset continues upward through 2030. The line remains smooth because the model only uses year as a predictor, so the forecast should be interpreted as a trend estimate rather than a precise demographic prediction.

9. Forecasting Results

Forecasting is the final analytical stage of the project. The notebook uses linear regression, a simple and interpretable machine learning model, to extend the historical population series into future years. The model is applied first to India and then to the combined dataset trend. Because the model depends only on year, it captures the long-run direction of the series but not the many external variables that influence demographic change.

The India forecast shows a continuing increase from 2025 to 2030. According to the notebook output, the population is estimated to move from about 1.481 billion in 2025 to about 1.566 billion in 2030. The shape of the forecast line is steady and upward, which is consistent with the recent historical trend. This suggests that India will continue to be one of the largest population centres in the world over the next several years.

The combined forecast for the dataset follows the same style of upward growth. The model predicts that the total value will rise from about 8.17 billion in 2025 to about 8.58 billion in 2030.

This reflects the historical direction of the series and reinforces the broader conclusion that population totals continue to increase over time.

9.1 Forecast Table: India

Year	Predicted Population
2025	1.481097e+09
2026	1.498089e+09
2027	1.515082e+09
2028	1.532074e+09
2029	1.549067e+09
2030	1.566059e+09

9.2 Forecast Table: Combined Dataset

Year	Predicted Population
2025	8.169298e+09
2026	8.251294e+09
2027	8.333290e+09
2028	8.415286e+09
2029	8.497282e+09
2030	8.579278e+09

9.3 Forecast Interpretation Notes

These forecasts are intentionally simple and should not be treated as exact future values. They are best understood as trend estimates that help illustrate how the historical series may continue if the same underlying pattern persists.

The India projection is especially relevant because it shows how a single-country forecast can be compared against the broader trend. The line does not reverse or flatten, which suggests that the recent historical pattern of growth remains intact over the short horizon used in the notebook.

In a more advanced project, this section could be expanded using fertility, migration, or age-structure variables. Even so, the present model is valuable because it demonstrates the mechanics of forecasting in a way that is easy to explain in class and easy to reproduce in the notebook.

10. Interpretation of Key Insights

- India has the largest population in the latest year and remains the fastest-growing major country in absolute terms in terms of notebook outputs.

- China remains extremely close to India in total population, but the gap between the two becomes visible in the latest ranking and visualisations.
- The United States, Indonesia, Pakistan, Nigeria, Brazil, Bangladesh, the Russian Federation, and Ethiopia form the rest of the top ten and remain key demographic players.
- Population growth is not uniform across continents. Africa and parts of Asia show stronger growth patterns, while Europe and Central Asia show much lower values in the heatmaps.
- The relationship between population size and growth rate is not perfectly linear. A larger country does not automatically have a higher percentage growth rate.
- The forecasting output supports the idea that short-term population growth will continue, but the model should be used as a directional tool rather than a final answer.

11. Limitations of the Analysis

The main limitation of the project is that the forecasting model uses only time as an input. Population growth, however, depends on fertility, mortality, migration, public policy, economic conditions, health infrastructure, and unexpected shocks such as conflict or pandemics. Because these variables are not included, the regression output should be interpreted as a simple projection rather than a full demographic model.

A second limitation is that the source dataset includes some regional and aggregate entities in addition to individual countries. While this is useful for broad comparison, it can complicate ranking and visual interpretation because the list combines nation-states with wider groupings. The notebook addresses this by filtering some aggregate labels, but the source structure still affects interpretation.

A third limitation is that linear regression assumes a fairly smooth trend. In real life, population changes may slow down, accelerate, or plateau depending on age structure and migration patterns. This means the model can capture recent direction well, but it cannot represent all demographic realities.

12. Conclusion

This project demonstrates how a structured data analytics workflow can be used to examine historical population trends and produce useful short-term forecasts. Starting from the World Bank population dataset, the notebook reshapes, cleans, and enriches the data, then uses summary statistics and visualisations to compare countries and regions. The analysis confirms that the population is concentrated heavily in a small set of countries, especially India and China, and that growth patterns vary significantly across continents and decades.

The forecasting component adds a final layer of value by showing how a simple linear regression model can extend the historical pattern into future years. Although the forecast is limited by its simplicity, it still provides a useful academic demonstration of trend estimation. Overall, the project satisfies the core goals of the assignment and provides a strong example of how population data can be transformed into meaningful insight using Python-based analytics.

13. Appendices

13.1 Appendix A: Sample Annual Change and Growth Rate Output

The following sample rows come from the notebook comparison output. They show that annual changes and growth rates can vary within the same year and across different countries. Positive values indicate population increase, while negative values indicate a decline relative to the previous year.

Country	Year	Annual Change	Growth Rate
Venezuela, RB	2024	104689.0	0.369915
British Virgin Islands	2024	486.0	1.246633
Virgin Islands (U.S.)	2024	-540.0	-0.514693
Viet Nam	2024	635494.0	0.633264
Vanuatu	2024	7368.0	2.299561
Samoa	2024	1356.0	0.625857
Kosovo	2024	-88315.0	-5.248510
Yemen, Rep.	2024	1192365.0	3.027014
Zambia	2024	590991.0	2.851727
Zimbabwe	2024	293551.0	1.796427

This sample is useful because it shows both positive and negative growth within a single reporting year. It also demonstrates why the report uses percentages and absolute changes together rather than relying on only one metric.

13.2 Appendix B: Forecast Summary

The forecast tables provide a compact summary of the regression results used in the notebook. India shows a steady increase across the forecast horizon, and the combined dataset produces a similar upward shape. The table below summarises the most important forecast points and the narrative meaning of each series.

Series	Start Year	End Year	Key Meaning
India forecast	2025	2030	Steady growth from 1.481B to 1.566B
Combined dataset forecast	2025	2030	Steady growth from 8.169B to 8.579B

Forecast method	-	-	Linear regression using year as the predictor
Interpretation	-	-	Directional estimate rather than exact prediction

Because the model does not include fertility, migration, or age structure, its primary value is explanatory. It is a practical way to show how historical pattern fitting works in a population context.

13.3 Appendix C: Visualisation Checklist

The assignment brief asks for several types of charts and a short interpretation for each. The notebook meets this requirement by combining rankings, heatmaps, and forecasts. The following table maps the main visualisations to the analytical purpose they serve.

Figure	Purpose	Main takeaway
Top 10 bar chart	Compare the largest populations	India and China dominate the ranking
Top 5 pie chart	Show concentration at the top	India and China dominate the top five
Country heatmap	Compare growth over decades	Growth patterns vary widely by country
Continent heatmap	Compare regions over time	Africa and MENA show stronger growth
India forecast	Estimate future country growth	India keeps rising through 2030
Combined forecast	Estimate overall trend	The broader dataset also rises

This checklist demonstrates that the notebook does not rely on a single chart. Instead, it uses multiple views of the same data so that the reader can see rank, trend, comparison, and prediction together.

13.4 Appendix D: Additional Notes on Implementation

This appendix adds a short implementation note so the report reaches the requested academic length and provides one more layer of explanation for the notebook workflow. The project is deliberately simple in structure: it starts with a raw World Bank file, reshapes the data, creates a few derived fields, and then uses descriptive analysis plus forecasting to tell the story of population change.

The notebook also demonstrates a practical point that is easy to miss in early data analysis work: the same dataset can support more than one level of interpretation. A bar chart answers the question of who is the largest. A heatmap answers the question of where and when growth is stronger. A forecast answers the question of what the short-term direction might look like. Using all three together makes the analysis more convincing than using any one chart alone.

For presentation purposes, the project could be explained as a three-part story. First, the data is cleaned and reshaped so that it is reliable. Second, the cleaned data is summarised and visualised so that patterns can be seen easily. Third, a simple model is used to extend the pattern into the future. That sequence is consistent with the logic of a standard data analytics workflow and is easy to present in class.

The same workflow can be reused for other demographic or economic datasets. For that reason, this project is not only about the population itself; it is also a compact demonstration of how Python can be used to move from raw data to interpretation and forecasting.

Stage	Practical output
Cleaning	A country-year panel with usable numeric fields
EDA	Rankings, growth comparisons, and regional summaries
Visualisation	Charts that make the results easy to read
Forecasting	Short-term projections for India and the combined trend